H3 experiments: foundation gains; adaptation shows a quality tradeoff

All displayed methods use the same output observation-consistency projection (DC25).

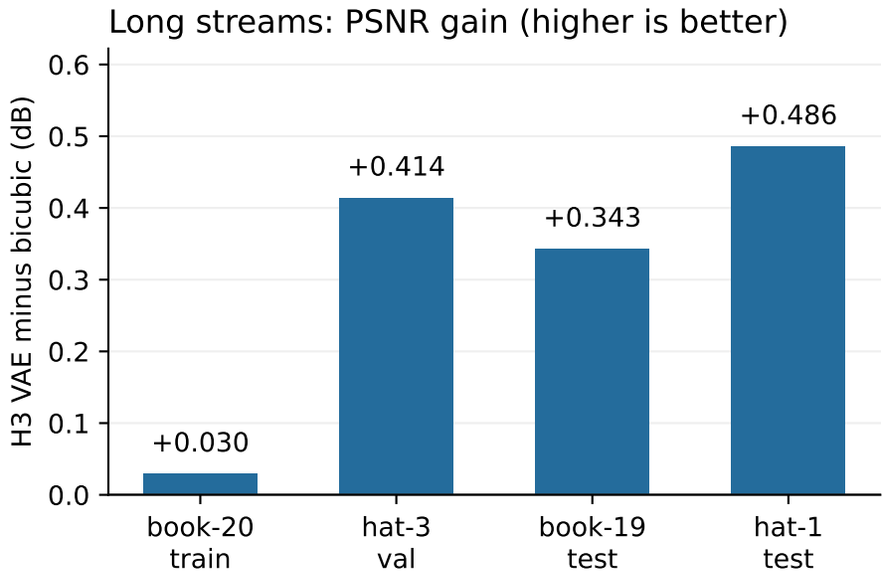

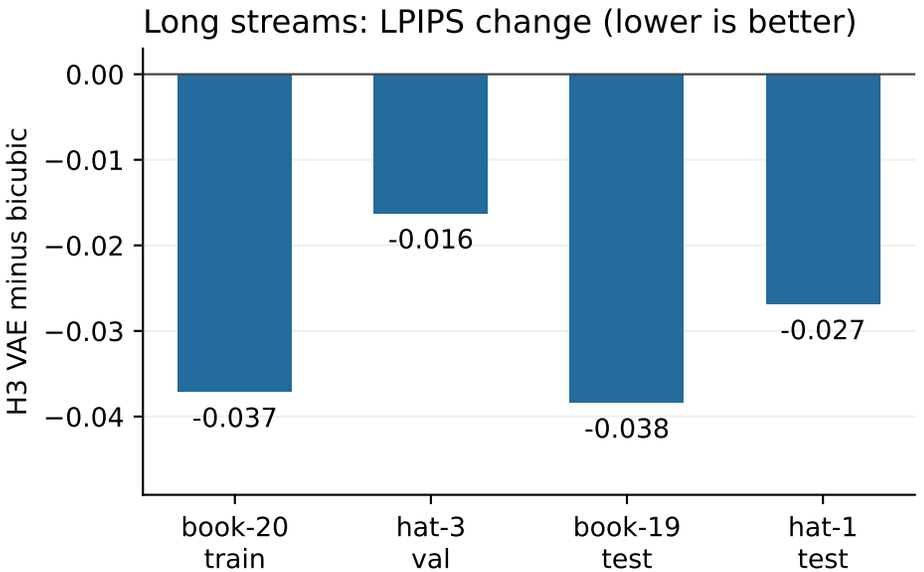

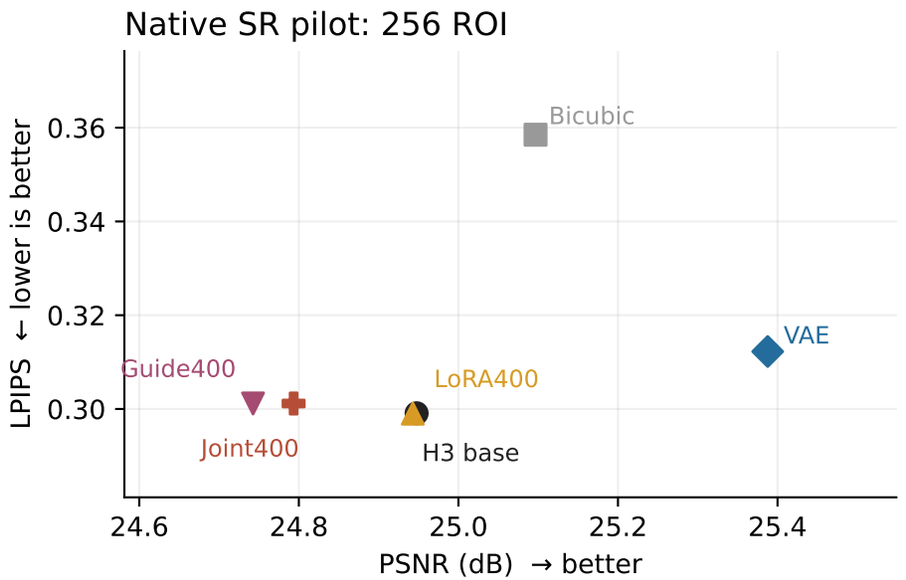

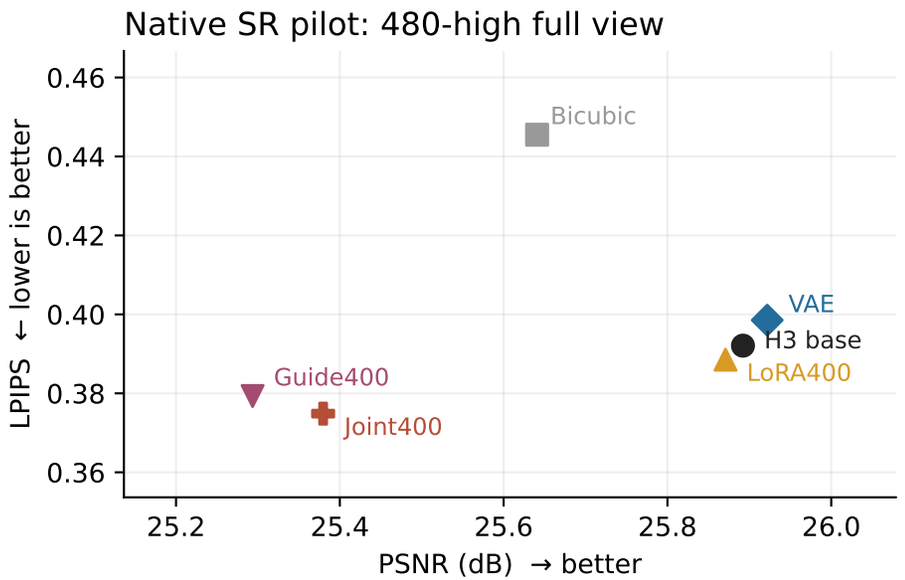

Top: 4 independent sources × 600 frames; frozen VAE, no persistent memory.
Bottom: v2 pilot, 4 training / 2 validation sources, 400 updates per control; training IO matches the official BF16 inference wrapper.
Spatial views differ; not a resolution ablation. This pilot does not establish an independent learned-memory benefit.